


Module 3 - Foundations of computer vision system (2) - Local feature and representation, part 2

Dr. Tan Jen Hong
Lecturer & Consultant
Institute of System Science
National University of Singapore
issjht@nus.edu.sg

Learning objectives


- Rescale / resize image
- Perform image translation and rotation
- Perform histogram equalization
- Perform image thresholding

What is interpolation?




Source: https://en.wikipedia.org/wiki/Interpolation#/media/File:Interpolation_Data.svg


Types of interpolation




1D nearest-neighbour




Linear




Cubic



2D nearest-neighbour







Bilinear



Bicubic

Source: https://en.wikipedia.org/wiki/Bicubic_interpolation


Types of interpolation



Source: https://en.wikipedia.org/wiki/Bicubic_interpolation


Roles of interpolation in image processing?

- Whenever we perform a geometric transformation on image, interpolation is always involved
- The choice of interpolation method affects the output and the calculation speed



Source: https://en.wikipedia.org/wiki/Bicubic_interpolation

Rescale image



cc.png

- Load the image and the necessary libraries

```
> import cv2  
> import numpy as np  
> import matplotlib.pyplot as plt
```

```
> src = cv2.imread('cc.png')
```

- Check the shape and the type of the image

```
> src.shape  
: (501, 501, 3)
```

```
> src.dtype  
: dtype('uint8')
```

```
> print('img height: %d' % (src.shape[0]))  
: img height: 501
```



Rescale image

- Two ways to rescale image. First method: specify the output size

```
> rsz1 = cv2.resize(src,      width height  
                    output size  dsize=(300,200),  
                    interpolation method  interpolation=cv2.INTER_AREA)
```

- Check the shape of the image

```
> rsz1.shape  
: (200, 300, 3)  
  
> print('img height: %d' % (rsz1.shape[0]))  
: img height: 200
```




rsz1

Rescale image

- Second method: specify scale factor in x and y directions

```
> rsz2 = cv2.resize(src,  
                    output size dsize=None,  
                    scale factor for width fx=0.5,  
                    scale factor for height fy=0.7,  
                    interpolation method interpolation=cv2.INTER_LINEAR)
```



- `dsize` must be set to `None` to avoid error
- Check the shape of the image

```
> rsz2.shape  
: (351, 250, 3)
```

```
> print('img height: %d' % (rsz2.shape[0]))  
: img height: 351
```

rsz2

Interpolation methods

- Opencv provides a number of interpolation methods for geometric transformation

- Some of the common used methods

```
cv2.INTER_AREA  
cv2.INTER_LINEAR  
cv2.INTER_CUBIC  
cv2.INTER_NEAREST
```


- ‘NEAREST’ has the fastest operation, but output is crude
- ‘AREA’ is good for shrinking image; ‘CUBIC’ good for enlarging image. But both operations are slow
- ‘LINEAR’ is somewhere in between

```
cv2.resize(src,  
           dsize,  
           interpolation=cv2.INTER_AREA)
```

Enlarging image


- Enlarging image rarely give satisfying output; avoid if possible
- Plan your processing flow carefully; keep the original source in your flow if possible

```
> sml = cv2.resize(src,  
                    dsiz=(150,150),  
                    interpolation=cv2.INTER_LINEAR)  
  
> bg1 = cv2.resize(sml,  
                    dsiz=(500,500),  
                    interpolation=cv2.INTER_NEAREST)  
  
> bg2 = cv2.resize(sml,  
                    dsiz=(500,500),  
                    interpolation=cv2.INTER_CUBIC)
```




sml

Enlarging image



nearest




cubic

Affine transformation

- Mathematically, scaling, translating, rotating an image are considered as affine transformation
- These transformations is ruled by a single matrix


Translate

$$\begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$



Scale


$$\begin{bmatrix} 2 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



Rotate


$$\begin{bmatrix} c & s & 0 \\ -s & c & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$c = s = \sin(45^\circ)$$



Shear

$$\begin{bmatrix} 1 & 0.5 & 0 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$



Source: <https://neutrium.net/mathematics/basics-of-affine-transformation/>


Affine transformation

- In the matrix the last row is always having the values $[0, 0, 1]$ for image processing purpose
- Thus in opencv, the last row is removed, and the ruling matrix is of this form

$$M = \begin{bmatrix} a_{11} & a_{12} & a_{tx} \\ a_{21} & a_{22} & a_{ty} \end{bmatrix}$$


Translate

$$\begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$



Scale


$$\begin{bmatrix} 2 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



Rotate


$$\begin{bmatrix} c & s & 0 \\ -s & c & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$c = s = \sin(45^\circ)$$



Shear

$$\begin{bmatrix} 1 & 0.5 & 0 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$



- Except for translation, we generally do not manually calculate each value in the matrix
- We try to use function to get the required transformation matrix

Source: <https://neutrium.net/mathematics/basics-of-affine-transformation/>

Translation


- For the translation, the transformation matrix looks like this

$$M = \begin{bmatrix} 1 & 0 & a_{tx} \\ 0 & 1 & a_{ty} \end{bmatrix}$$

- This assumes only translation operation is applied; no scaling or rotation is performed.
- a_{tx} stands for the amount of distance for the image to be shifted along x axis (in pixels)
 - positive value shifts the image to the right, negative to the left
- Similarly, a_{ty} stands for the shift in y axis
 - positive value shifts the image to the bottom, negative to the top

Translate

$$\begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$



Source: <https://neutrium.net/mathematics/basics-of-affine-transformation/>

Translation

- To translate an image in opencv, the first thing we do is to create the transformation matrix

```
> Mt = np.float32([[1, 0, 0],  
                   [0, 1, 25]])
```


a_{tx}
a_{ty}

- Must use `np.float32` to create the transformation matrix, else error

- `Mt` shifts an image downward by 25 pixels

```
> tr = cv2.warpAffine(src,  
                       transformation matrix Mt,  
                       output size (501, 501))
```

height width



tr

Rotation

- To create the transformation matrix for rotation, we do

```
> (row, colm, _) = src.shape
```


```
> Mr = cv2.getRotationMatrix2D((colm/2, row/2),  
                                90, 1)
```

rotation center
rotation angle
scale

- Opencv provides the function to generate transformation matrix for rotation about a center point in an image
- The rotation angle is in degree, positive value for anti-clockwise, neative value for clockwise

```
> rt = cv2.warpAffine(src,  
                       transformation matrix Mr,  
                       output size (501, 501))
```

height width



A very important use case

On geometric transformation




Source: <https://rock-it.pl/images-augmentation-for-deep-learning-with-keras/>

Image histogram

Image histogram



- Image histogram: a graph of pixel intensity (along x-axis) vs number of pixels (y-axis)
- x-axis has all the gray levels
- y-axis indicates the amount of pixels in image that have the particular gray-level value



Source: <https://www.allaboutcircuits.com/technical-articles/image-histogram-characteristics-machine-learning-image-processing/>

Image histogram



- 4 x 4 monochrome image (black or white) image
- Count the number of black pixels (there are 4)
- Count the number of white pixels (there are 12)



Source: <https://www.allaboutcircuits.com/technical-articles/image-histogram-characteristics-machine-learning-image-processing/>

Image histogram





- 4 x 4 gray scale image
- Gray scale image are generally 8-bit, and thus it has 255 levels
- Note: in image processing, we sometimes group several levels together as one bin



Source: <https://www.allaboutcircuits.com/technical-articles/image-histogram-characteristics-machine-learning-image-processing/>

Image histogram


- Problem with histogram: provides only statistical information
- Zero information about the spatial distribution of the pixel values



Source: <https://www.allaboutcircuits.com/technical-articles/image-histogram-characteristics-machine-learning-image-processing/>

Image histogram


- Possible solutions: divide image into regions and do further analysis



Source: <https://towardsdatascience.com/face-recognition-how-lbph-works-90ec258c3d6b>

Image histogram

- Image of high contrast and low contrast



Source: <https://www.allaboutcircuits.com/technical-articles/image-histogram-characteristics-machine-learning-image-processing/>

Image histogram



kwg.jpg

- Let us load an image

```
> kwg = cv2.imread('kwg.jpg')
```

- Check the shape

```
> kwg.shape  
: (700, 477, 3)
```

- Grayscale image by right, should have only one channel
- However, many times grayscale images are still saved in 3 channels
- To retrieve just one channel for processing, we do

```
> kwg = kwg[:, :, 0]
```


Image histogram

- Plot the image histogram

```
> plt.hist(kwg.ravel(),  
           number of bins 256,  
           range [0,255],  
           colour of the bar facecolor='black')
```




kwg.jpg




Histogram equalization

- Histogram equalization: a method to adjust contrast in an image
- Achieved by spreading out most frequent intensity values
- Often produce unrealistic effects in photographs, but useful for scientific images
- Calculation is not computationally expensive




Source: <https://commons.wikimedia.org/wiki/File:Histogrammspreizung.png>

Why do we need to adjust contrast?



Source: https://www.freepik.com/free-icon/brightness-and-contrast-adjustment-option_748892.htm




Source: https://www.freepik.com/free-icon/contrast-adjustment-symbol_736518.htm


Histogram equalization

- To perform histogram equalization, simply do


```
> heq = cv2.equalizeHist(kwg)
> plt.hist(heq.ravel(),
            number of bins 256,
            range [0,255],
            colour of the bar facecolor='black')
```



heq



kwg




Contrast limited adaptive histogram equalization (CLAHE)

- To perform CLAHE, we do


```
> clahe = cv2.createCLAHE()  create CLAHE object
> hcl    = clahe.apply(kwg)   use the object to
> plt.hist(hcl.ravel(),      perform equalization
                        256,
                        [0,255],
                        facecolor='black')
```



hcl



kwg



Comparison

Histogram equalization




CLAHE



Image thresholding

- Image thresholding: a process that converts intensity values that exceeds/below a threshold to one value, and the rest as other value(s)
- Usually we use 1 or 255 to denote the pixels/points that we want, and 0 for points/pixels to be ignored/discarded
- This is often used to isolate region of interest



Source: <http://www.ni.com/tutorial/2916/en/>


Image thresholding

- When we do thresholding, we are in fact performing segmentation using the image histogram

original image



thresholded image



Source: <https://www.allaboutcircuits.com/technical-articles/image-histogram-characteristics-machine-learning-image-processing/>

Image thresholding

- Opencv provides several methods to do thresholding

```
> gr = cv2.imread('gr.png')  
> (_,thr) = cv2.threshold(gr[:, :, 0],  
                           threshold value 85,  
                           assigned value 255,  
                           threshold style cv2.THRESH_BINARY)
```

- Thresholding can be easily implemented in simple python; just one/two for loops will do
- But that will be slow; using opencv function gives faster implementation




gr.png


Thresholding style

```
cv2.threshold(gr[:, :, 0],  
             85,  
             255,  
             ThresStyle)
```


GR.PNG




BINARY




BINARY_INV



TRUNC



TOZERO



TOZERO_INV

