

# **Module 4 - Foundations of computer vision system (3) - Global feature and representation, part 1**

Dr. Tan Jen Hong  
Lecturer & Consultant  
Institute of System Science  
National University of Singapore  
[issjht@nus.edu.sg](mailto:issjht@nus.edu.sg)

# Learning objectives

- Perform image filtering
- Perform edge detection
- Extract objects from image
- Draw contour around objects

# Undesirable effect from uint8

- Let's investigate uint8, declare two variables, assumed already import numpy as np

```
> A = np.uint8(5)
> B = np.uint8(8)
> C = A - B
```

- What is the output of C?
- The way numpy / python handles the above situation is problematic in image processing
- Thus, many times we need to convert uint8 to float32

## Handling something beyond boundaries ...

- Many times, the values of the output of a processing exceeds the usual range
- The values either exceed 255 or go into negative (for integer)
- The values either exceed 1 or below 0 (for float)
- There is a need to rescale the values back to the desired range
- We call this process normalization

## Handling something beyond boundaries ...

- Let  $x$  denote the pixel values in an image; let oldMin and oldMax denote the minimum and the maximum values respectively in the image
- let  $x_n$  denote the normalized value; let newMin and newMax denote the desired minimum and maximum values
- The formula to rescale / re-normalize:

$$\frac{x - \text{oldMin}}{\text{oldMax} - \text{oldMin}} = \frac{x_n - \text{newMin}}{\text{newMax} - \text{newMin}}$$

## Handling something beyond boundaries ...

- Re-arrange and we get:

$$x_n = \frac{x - \text{oldMin}}{\text{oldMax} - \text{oldMin}} (\text{newMax} - \text{newMin}) + \text{newMin}$$

- Exercise: Create a function with the below signature to normalize an image (the `img` must be a float)

```
> def imgNormalize(img, minv=0, maxv=1):
```

# Image filtering

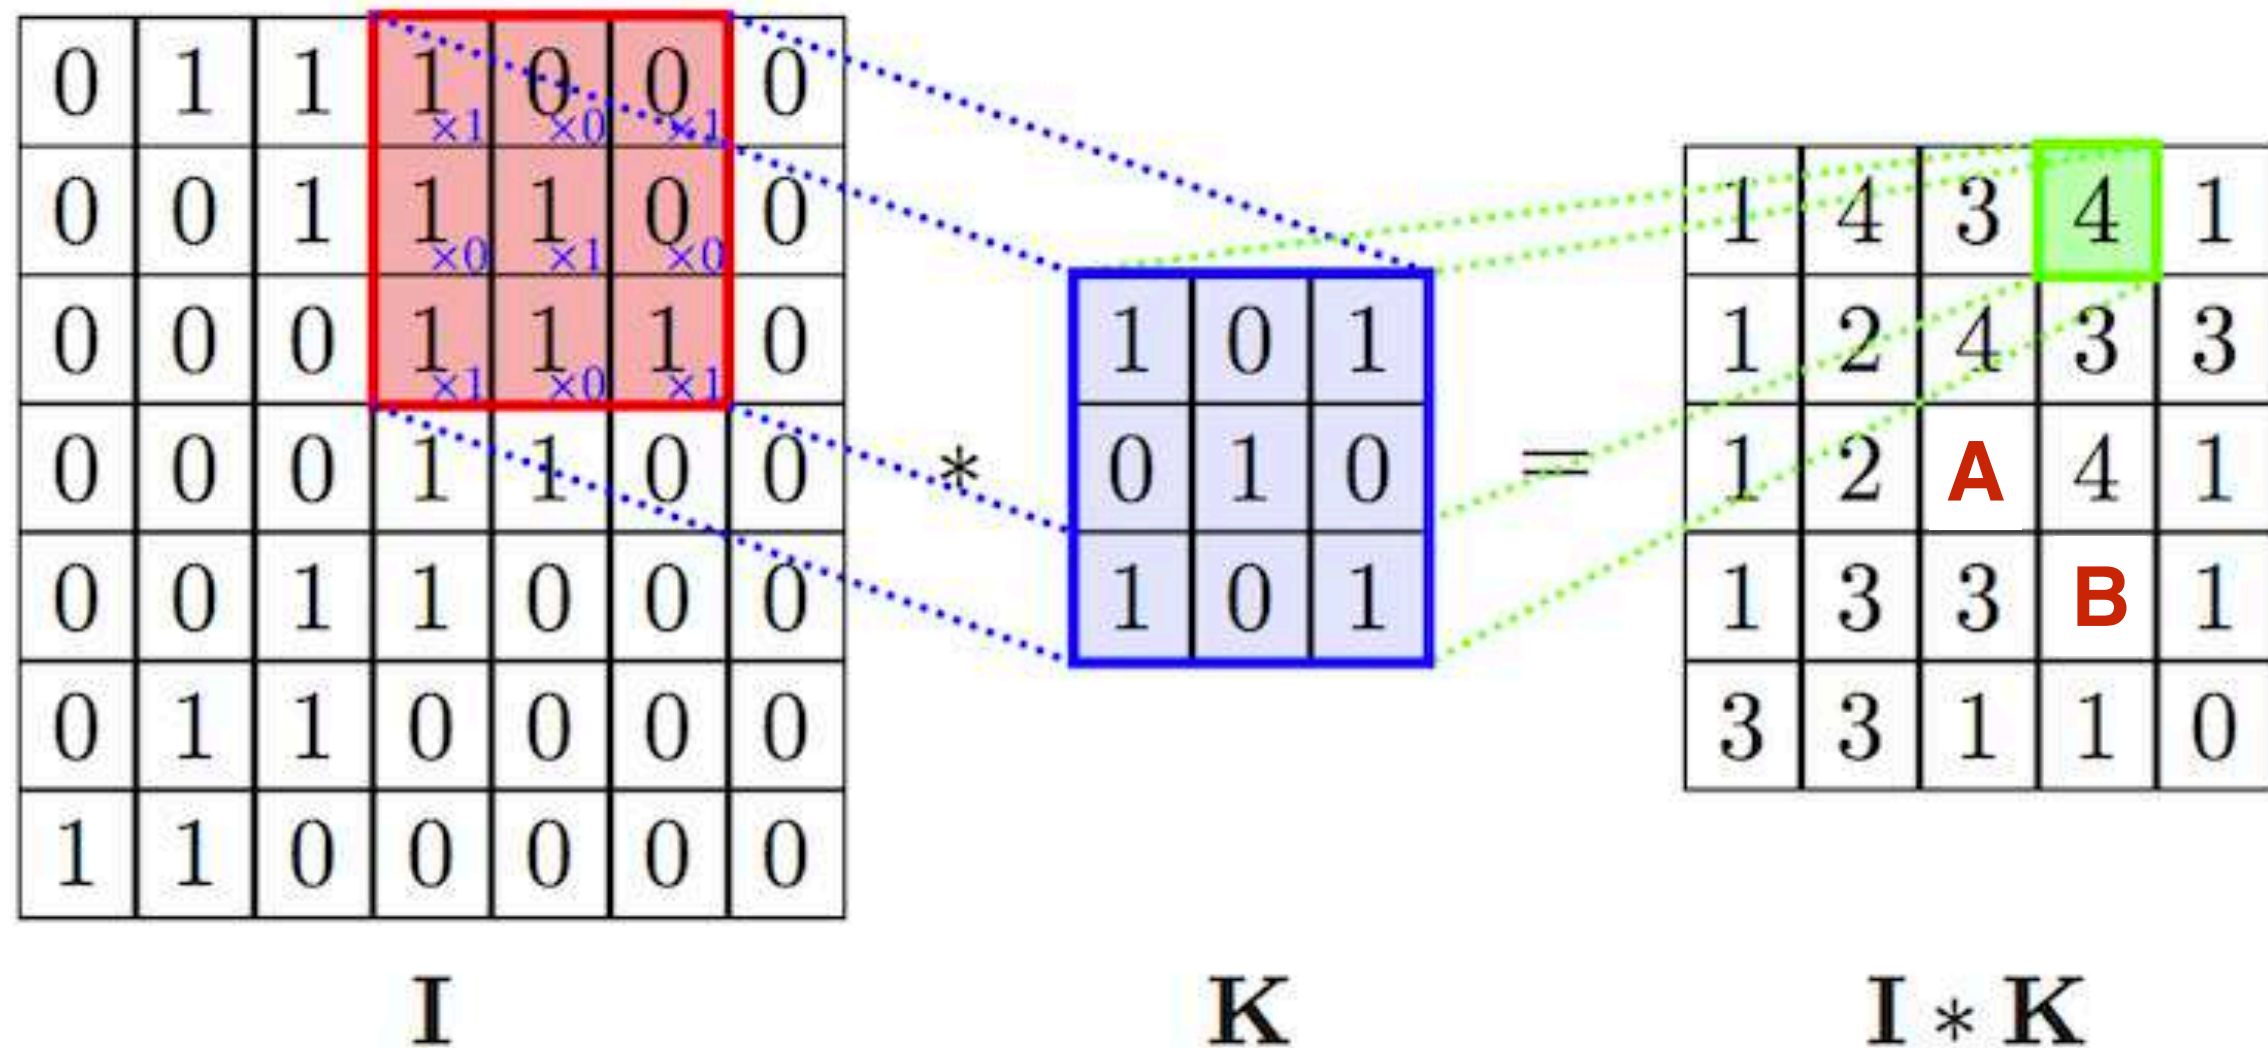
- In image processing, filtering is used to strengthen preferred features or weaken unwanted features
- We use filtering to smooth or sharpen an image; also, to enhance edges in image
- Filtering is achieved through convolution; it involves a kernel and the image of interest
- The design of the kernel determines the filtering output



Source: <https://www.star-spain.com/en/blog/transittermstar-nxt-tooltips/filtering-data-records-termstar-nxt>

# 2D Convolution

- What is the value of A and B?

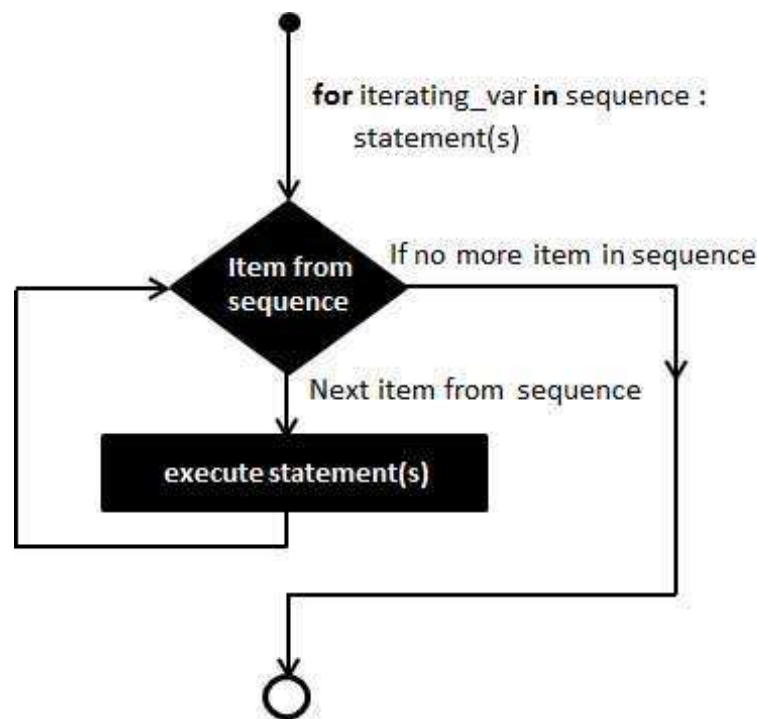


Source: [https://www.researchgate.net/figure/An-example-of-convolution-operation-in-2D-2\\_fig3\\_324165524](https://www.researchgate.net/figure/An-example-of-convolution-operation-in-2D-2_fig3_324165524)



# Note on 2D Convolution

- Avoid writing convolution code in python
- for loops in python are very slow; the calculation takes much longer time
- Always use library functions; they are optimized
- If not, write C code



Source: [https://www.tutorialspoint.com/python/python\\_for\\_loop.htm](https://www.tutorialspoint.com/python/python_for_loop.htm)

# Some convolutions



Original

|    |    |    |
|----|----|----|
| •0 | •0 | •0 |
| •0 | •1 | •0 |
| •0 | •0 | •0 |

=

?



Original

|    |    |    |
|----|----|----|
| •0 | •0 | •0 |
| •0 | •0 | •1 |
| •0 | •0 | •0 |

=

?



Original

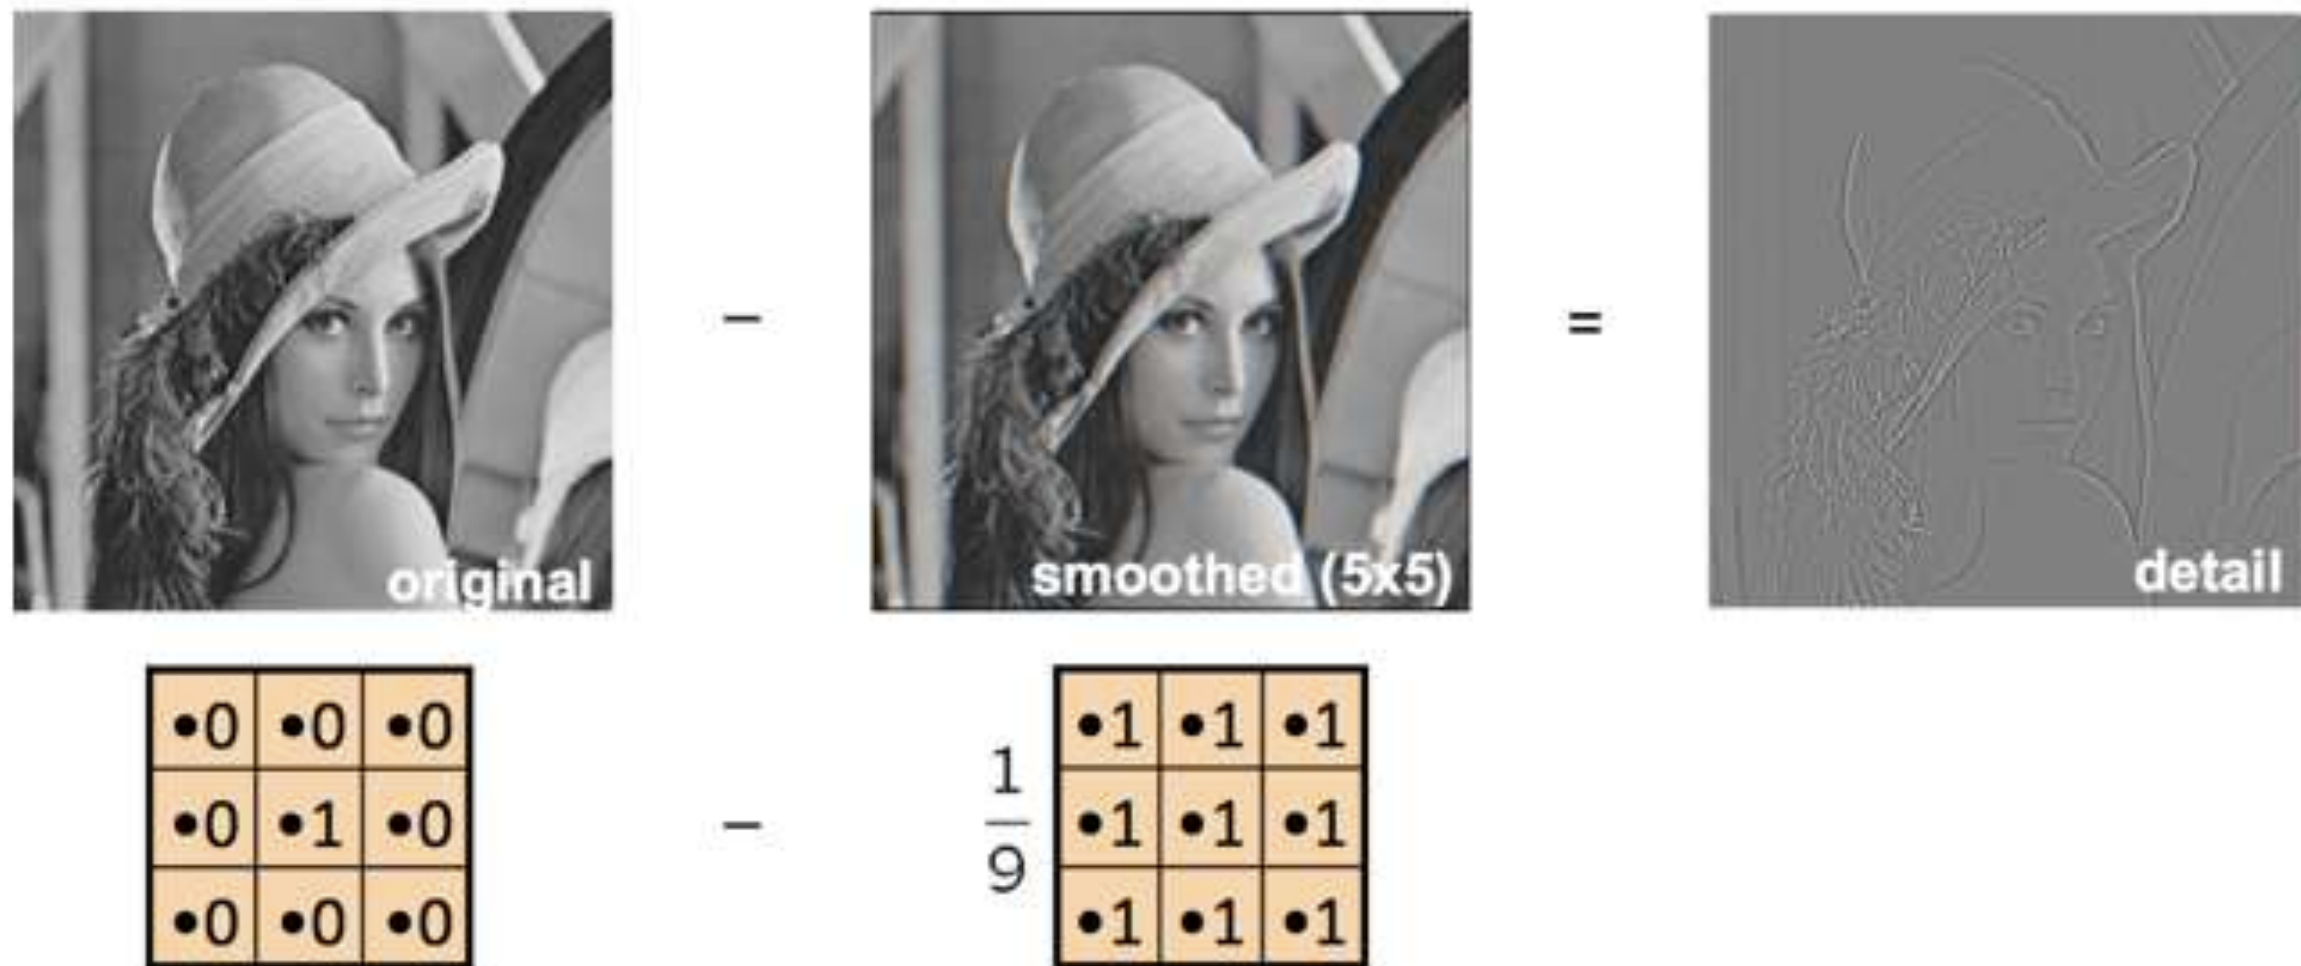
|               |    |    |    |
|---------------|----|----|----|
| $\frac{1}{9}$ | •1 | •1 | •1 |
|               | •1 | •1 | •1 |
|               | •1 | •1 | •1 |

=

?

Source: <https://ai.stanford.edu/~syYeung/cvweb/tutorial1.html>

# Some combinations



Source: <https://ai.stanford.edu/~syueung/cvweb/tutorial1.html>

# Some combinations



$$\begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix} + \begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 1 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix} = \begin{bmatrix} \bullet 0 & \bullet 0 & \bullet 0 \\ \bullet 0 & \bullet 2 & \bullet 0 \\ \bullet 0 & \bullet 0 & \bullet 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix}$$

Source: <https://ai.stanford.edu/~syYeung/cvweb/tutorial1.html>

# Basic filtering

- Load the image and the necessary libraries

```
> import cv2
> import numpy as np
> import matplotlib.pyplot as plt

> dr3 = cv2.imread('dr3.png')
```

- Create the kernel

```
> knl = np.ones((7,7),
                 np.float32)/49
```

- Do the filtering

```
> bsc = cv2.filter2D(dr3,
                      intensity depth, just set to -1 -1,
                      kernel knl)
```




dr3.png



# Basic filtering

Original



Filtered



# Mean filtering

- Mean filter reduces intensity variation among pixels
- Commonly used to reduce noise in images
- For a kernel of size  $j \times j$ , each entry in the matrix shares the same value of  $1/(j \times j)$
- Example: a  $5 \times 5$  kernel for a mean filter looks like

$$\frac{M + E + A + N}{4}$$

$$\begin{bmatrix} \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \\ \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \\ \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \\ \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \\ \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \end{bmatrix}$$

Source: <http://www.webquest.hawaii.edu/kahihi/mathdictionary/M/mean.php>


# Mean filtering

- In opencv, mean filtering is simply by


```
> nuc = cv2.imread('nuc.jpg')  
> mflt = cv2.blur(nuc, (11, 11))
```

kernel  
size

nuc



mflt






# Gaussian filtering

- Gaussian filter: non-uniform low pass filter
- Commonly used to remove noise and detail
- The formula for 2D kernel (isotropic):

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2 + y^2}{2\sigma^2}}$$



- When we use Gaussian filter, sigma is the key
- The kernel size should be around 6 x sigma (3 sigma for each side)


Source: <https://www.teamliquid.net/forum/starcraft-2/142211-sc2-ladder-analysis-part-2>

# Gaussian filtering


- In opencv, Gaussian filtering is simply by

```
> gaus= cv2.GaussianBlur(nuc,  
                           kernel size (11,11),  
                           sigmaX, set to 0, so that  
                           opencv will calculate the  
                           sigmaX and sigmaY based on  
                           the kernel size  
                           0)
```

nuc



gaus



# Median filtering

- Median filter: it considers the neighbourhood around a pixel, order the values, determine the median value and output the median value
- Commonly used to remove noise and detail; do better job than mean filter

|     |     |     |     |     |
|-----|-----|-----|-----|-----|
| 123 | 125 | 126 | 130 | 140 |
| 122 | 124 | 126 | 127 | 135 |
| 118 | 120 | 150 | 125 | 134 |
| 119 | 115 | 119 | 123 | 133 |
| 111 | 116 | 110 | 120 | 130 |

**Neighbourhood values:**

115, 119, 120, 123, 124,  
125, 126, 127, 150

**Median value: 124**

Source: <https://homepages.inf.ed.ac.uk/rbf/HIPR2/median.htm>


# Median filtering

- In opencv, Median filtering is simply by


```
> medf= cv2.medianBlur(nuc,11)
```

kernel  
size

nuc




medf







nuc




mflt



gauss



medf



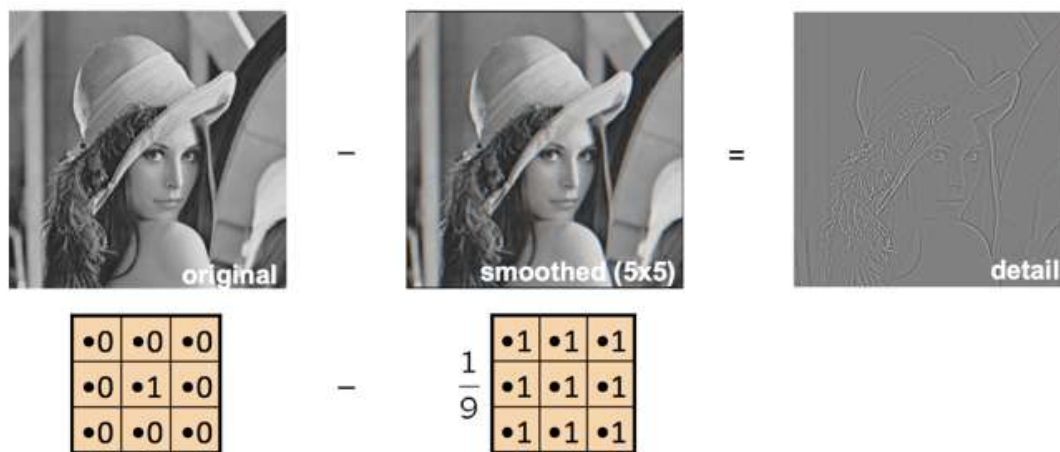
# How about sharpen the image?

- Create the mean filter kernel and perform the mean filtering

```
> krn = np.ones((11,11),np.float32)/121  
> mnuc = cv2.filter2D(nuc,-1,krn)
```

- Perform the subtraction and plot the output

```
> detail = (np.float32(nuc) - np.float32(mnuc))/255  
> cv2plt(detail)
```



Source: <https://ai.stanford.edu/~syeyeung/cvweb/tutorial1.html>

## How about sharpen the image?

- Can't see anything .....
- Let's check the maximum and minimum value
- Max: 0.29; Min: -0.27





# How about sharpen the image?

- Rescale the image, use the normalization function

```
def imgNormalize(img, minv=0, maxv=1):
```

- The code to achieve that:

```
> normalizedDetail = imgNormalize(detail)
> plt.figure()
> cv2plt(normalizedDetail)
```





# Sharpen the nuc



- To sharpen the image, we do

```
> sharpen = (np.float32(nuc)*2 - np.float32(mnuc))/255
```

- And plot the output:


```
> plt.figure()
```

```
> cv2plt(normalizedDetail)
```



## Before and after


- The effect of sharpening: Sharpen the boundaries but also introduce noise in the image



# Edge detection

# Edges and contours



- Local changes in intensity or colour, such as edges and contours, are important for visual perception
- In human visual system, a few lines in caricature and illustration are often sufficient to describe an object and a scene
- This shows how much edges and contours can inform with our visual system



Source: Digital Image Processing by Burger and Burge, 2016

# Edges and contours

- Edges are usually located at areas where the change in intensity values is large




Source: Digital Image Processing by Burger and Burge, 2016



# Edges and contours

- Edges are usually located at sites where the change in intensity values is large, i.e. the gradient at the site is large
- Reality: Not all the large changes in intensity values are the edges we are interested in
- Key to good edge detection: get only the edges we want, and eliminate the rest



Source: Digital Image Processing by Burger and Burge, 2016

# Edge detection

- Strategies to edge detection: Calculate the gradient in x-direction and y-direction
- The calculation of gradient generally involves 2D convolution with specific kernels. Assume  $I$  is the image, and the kernel is  $H_x$  and  $H_y$


$$I_x = I * H_x \quad I_y = I * H_y$$

- $I_x$  and  $I_y$  is the gradient in x and y direction respectively, and the local edge strength is given by

$$E(i, j) = \sqrt{I_x^2(i, j) + I_y^2(i, j)}$$

- The local edge orientation angle is given by


$$\Phi(i, j) = \tan^{-1} \left( \frac{I_y(i, j)}{I_x(i, j)} \right)$$



Source: Digital Image Processing by Burger and Burge, 2016

# Edge detection

- The flow of getting the edges are as below



Source: Digital Image Processing by Burger and Burge, 2016



# Edge detection

- Some of the common kernels / operators to get edges are Prewitt, Roberts and Sobel operators
- Roberts operators are one of the simplest and oldest methods; Prewitt and Sobel are almost similar
- The kernels for Sobel operator:



$$H_x^S = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} \quad \text{and} \quad H_y^S = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

# Sobel edge operators


original image



local edge strength




local edge orientation  
(The edge orientation estimated by Sobel is relatively inaccurate)



Source: Digital Image Processing by Burger and Burge, 2016

## After gradient, what's next?

- Edge operator generally generates values that indicate the likelihood of a pixel to be part of edge
- The problem: How to judge if a pixel is part of edge?
- Use of thresholding? This strategy often produces many unconnected edge segments
- Tracing edges pixel by pixel? This strategy does not work well in scenario where edges cross, or a single edge branches into several direction
- Solution? Canny edge detection can help, but still can't fully solve the problems




Source: Digital Image Processing by Burger and Burge, 2016

# Canny edge detection


- There are four steps in Canny edge detection
- 1. Noise reduction: remove noise with 5 x 5 Gaussian filter
- 2. Calculate intensity gradient: Filter image with Sobel kernel in both horizontal and vertical direction. In image processing, Sobel kernel is often used to calculate the range of change in pixels along x-direction and y-direction
- 3. Non-maximum suppression: check if an edge point is a local maximum in its neighbourhood in the direction of gradient
- 4. Hysteresis thresholding: everything above maxVal is edge; every thing below minVal is not edge; For those in between, see connections



# Canny edge detection



Sobel operator



Canny edge detection


Source: Digital Image Processing by Burger and Burge, 2016

# Canny edge detection


- In opencv, canny edge detection is simply done by

```
> cann= cv2.Canny(nuc,  
                  minVal 31,  
                  maxVal 127,  
                  Sobel kernel size apertureSize=3)
```

nuc




cann




# Canny edge detection

- The effect of different maxVal

maxVal = 127




maxVal = 191






# Problem: How to isolate region of interest?

neu.jpg



wks2\_2\_e.jpg





# Region isolation

- Load the images


```
> neu = cv2.imread('neu.jpg')  
> msk = cv2.imread('wks2_2_e.jpg')
```

- Process the mask

```
> msk = np.float32(msk[:, :, 0])/255  
> msk = np.round(msk)
```

- Do the isolation

```
> iso = cv2.bitwise_and(neu,   
                           source 1 neu,   
                           source 2 neu,   
                           mask mask=np.uint8(msk))
```



# Draw boundaries

- To draw boundaries, first go and find boundary contour

```
> ctrs= cv2.findContours(np.uint8(msk), mask
                        contour retrieval modes cv2.RETR_EXTERNAL,
                        contour approximation methods cv2.CHAIN_APPROX_SIMPLE)
```

- `ctrs` is a tuple with 2 items. First item is a list of contour; second item is a hierarchy matrix

```
> ctrs= ctrs[0]
```

- Draw the contours

```
> cv2.drawContours(neu, image to be drawn
                  ctrs, list of contours
                  contour index; -1 indicate all contours to be drawn -1,
                  colour of the contours (191, 191, 255),
                  thickness 5)
```

# Draw contours

