

# **Module 5 - Building vision system using machine learning (1) - Detection and recognition, part 2**



Dr. Tan Jen Hong  
Lecturer & Consultant  
Institute of System Science  
National University of Singapore  
[issjht@nus.edu.sg](mailto:issjht@nus.edu.sg)

# Learning objectives

- Understand the difference between image classification and object detection
- Understand the major challenges in object detection
- Perform object detection using YOLO v3

# Image classification

- Algorithm looks at image and classifies the object within
- Meant for single item identification or category classification
- Wide application, from face detection on social network to medical diagnosis in healthcare



# Object classification and localization

- Tell not just if something exists in an image, but where it situates
- Draw bounding box and label



Source: <https://mangsh.co/2014/09/16/walking-waiting/>

# Object classification and localization

Before the deep learning era

- Before we introduce the deep learning approach to object classification and localization, let's see how this was done in the past
- Goal: Locate carplate in an image
- Using only the concepts we have covered in these 3 days

Before



After





# Locate carplate

## preprocessing

- Colour image (which is 3D) has rich information, but 2D array is easier for many image processing techniques

```
> car = cv2.imread('carplate01.png')  
> cargray = cv2.cvtColor(car, cv2.COLOR_BGR2GRAY)
```

Before



After



# Locate carplate

Apply thresholding

- Use Otsu thresholding to separate foreground and background in the scene

```
> lightest = cv2.threshold(cargray,
```

in Otsu thresholding, we set 0 for threshold value **0**,

assigned value **255**,

There are 2 outputs from Otsu thresholding, we are only interested in the second output `cv2.THRESH_OTSU)` **[1]**

Before



After



# Locate carplate

## Edge detection

- After thresholding, we need to extract the salient features in the image, this is done through edge detection

```
> edges = cv2.Canny(lightest,  
                    minVal 31,  
                    maxVal 127,  
                    Sobel kernel size apertureSize=3)
```

Before



After





# Locate carplate

Fuse characters on carplate

- To identify the carplate, we need those characters on the carplate to be fused as a cohesive block, to do that, we need to dilate and erode the borders of the characters

```
> bulk = cv2.dilate(edges, None, iterations=4)  
> bulk = cv2.erode(bulk, None, iterations=4)
```


Since we just use the default kernel (size 3 x 3), and thus we set 'None' for kernel

Each operation is run 4 times, hence the value 4 for 'iterations'

Before



After



# Locate carplate


Remove background noise

- To remove background noise, we apply erode followed by dilate 2 times

```
> bulk = cv2.erode(bulk, None, iterations=2)
> bulk = cv2.dilate(bulk, None, iterations=2)

> bulk = cv2.erode(bulk, None, iterations=4)
> bulk = cv2.dilate(bulk, None, iterations=4)
```

Before



After



# Locate carplate

Get the contours

- Get some of the largest contours in the background-remove image, and plot them out for inspection

Get all the contours in the image

```
> ctrs = cv2.findContours(bulk.copy(),  
                           cv2.RETR_EXTERNAL,  
                           cv2.CHAIN_APPROX_SIMPLE)
```

```
> ctrs = ctrs[0]
```

Sort the contours based on contour area, keep only the first 5 largest contour

```
> ctrs = sorted(ctrs,  
                key=cv2.contourArea,  
                reverse=True)[:5]
```

```
> cargrayrgb = cv2.merge((cargray, cargray, cargray))
```

```
> cv2.drawContours(cargrayrgb,  
                  ctrs,  
                  -1,  
                  (0, 0, 255),  
                  2)
```

Merge grayscale image into a 3D array, draw contours on the 3D array



# Locate carplate

Get the contours

- Only the first 5 largest-area contours are highlighted and kept in the analysis

Before



After





# Locate carplate

Get the carplate

- Apply bounding rectangle on each of the contour and determine the carplate by aspect ratio

```
> for c in ctrs:
    Apply bounding rectangle and calculate the aspect
    ratio for each contours
    (x, y, w, h) = cv2.boundingRect(c)
    aspr = w/float(h)
```

If the bounding box has an aspect ration between 4 and 5, it is considered the carplate

The carplate region is extracted and applied Otsu thresholding to get the car numbers

```
if aspr >= 4 and aspr <= 5:
    carPlate = cargray[y:y+h, x:x+w]
    carPlate = cv2.threshold(carPlate,
                             0,
                             255,
                             cv2.THRESH_OTSU) [1]

    break
```

```
> cv2.rectangle(car,
                (x,y),
                (x+w,y+h),
                (0,0,255),
                2)
```

Draw the carplate bounding box on the colour image



# Locate carplate

Get the contours

- The located carplate and the extracted region
- We can further send the extracted region to OCR library to read the characters

Before




After



# Multiple objects detection

- We have seen how in the past objects are located. But in the previous example, we are only locating one type of object
- Locate multiple objects (various categories) in an image is much more challenging
- Thus nowadays deep learning is the main solution to locate multiple objects in a scene
- Well known application: self-driving car; need to detect cars but also pedestrians



Source: <https://mangsh.co/2014/09/16/walking-waiting/>



# Why

- Why in deep learning can't we locate objects just sliding a box around image and performing classification on each extracted portion?





# Primary object detectors

- Faster R-CNN, by Girshick et al. 2015; <https://arxiv.org/abs/1506.01497>
- Single shot detectors (SSD), by Liu et al., 2015; <https://arxiv.org/abs/1512.02325>
- YOLO v3, by Redmon and Farhadi, 2018; <https://arxiv.org/abs/1804.02767>
- RetinaNet, by Lin et al., 2017; <https://arxiv.org/abs/1708.02002>

# YOLO v3

- YOLO stands for "you only look once", a one-stage detector
- One-stage detectors are faster, though less accurate compared to two-stage detector
- It is so fast to be used for tasks that require real-time response



Source: <https://pythonawesome.com/yolov3-training-and-inference-in-pytorch/>

# YOLO v3

- YOLO does detection at 3 stages / scales





# YOLO v3

- At first stage, YOLO divides image into  $13 \times 13$  cells
- In this stage YOLO tries to detect big objects in the image
- For each cell, the algorithm does object detection with three anchor boxes / prior boxes
- In total there are  $13 \times 13 \times 3 = 507$  anchor boxes in this stage





# YOLO v3

- For each anchor box, YOLO locates the object and return the bounding box that encloses the object
- It also makes prediction on the class of the object within the bounding box



# YOLO v3


- At second stage, YOLO divides image into  $26 \times 26$  cells
- In this stage YOLO tries to detect medium objects in the image
- For each cell, the algorithm does object detection with another three anchor boxes / prior boxes
- In total there are  $26 \times 26 \times 3 = 2028$  anchor boxes in this stage





# YOLO v3

- For each anchor box, YOLO locates the object and display the result with bounding box
- It then makes prediction on the class of the object within the bounding box




# YOLO v3

- At third stage, YOLO divides image into  $52 \times 52$  cells
- In this stage YOLO tries to detect small objects in the image
- For each cell, the algorithm does object detection with three small anchor boxes / prior boxes
- In total there are  $52 \times 52 \times 3 = 8112$  anchor boxes in this stage
- (Note: The anchor boxes are too small to be illustrated properly in the left image)





# COCO dataset



In total:  
123,287 images; 886,284 instances  
80 categories

Source: <http://cocodataset.org>

# YOLO Implementation

- Load the label file

```
> import cv2
> import numpy as np
> import matplotlib.pyplot as plt

> lbl_file = 'yolov3.txt'
> classes = open(lbl_file).read().strip().split("\n")
```

- Load the net model

```
> yoloconfig = 'yolov3.cfg'
> yoloweights= 'yolov3.weights'
> net          = cv2.dnn.readNet(yoloweights,
                                yoloconfig)
```

## yolov3.txt

```
person
bicycle
car
motorcycle
airplane
bus
train
truck
boat
traffic light
fire hydrant
stop sign
parking meter
bench
bird
cat
dog
horse
sheep
cow
elephant
bear
zebra
.....
```

## yolov3.cfg

```
[net]
# Testing
batch=1
subdivisions=1
# Training
# batch=64
# subdivisions=16
width=416
height=416
channels=3
momentum=0.9
decay=0.0005
angle=0
saturation = 1.5
exposure = 1.5
hue=.1

learning_rate=0.001
burn_in=1000
max_batches = 500200
policy=steps
steps=400000,450000
scales=.1,.1
.....
```

# YOLO Implementation

- Read image and create blob

```
> img = cv2.imread('ms3.jpg')
> blob = cv2.dnn.blobFromImage(image=img,
                                scaling for image values (not image size) scalefactor=1/255,
                                output size size=(416, 416),
                                mean value to be subtracted from each channel mean=(0,0,0),
                                swap R and B channel swapRB=True,
                                Do cropping after resizing image crop=False)
```

- Store the image height and width

```
> imgHeight = img.shape[0]
> imgWidth = img.shape[1]
```



Source: <https://mangsh.co/2014/11/09/of-street-01/>



# YOLO Implementation

- Create a function to get output layers since there are 3 output layers, each layer generates output at a particular scale

```
> def getOutputLayers(net):  
    layers = net.getLayerNames()  
    outLayers = [layers[i[0] - 1] for i in net.getUnconnectedOutLayers()]  
  
    return outLayers
```

- Set input, get output layers, run YOLO

```
> net.setInput(blob)  
> outLyrs = getOutputLayers(net)    get output layers  
> preds = net.forward(outLyrs)     Run the actual object detection.
```



The output is a list, consisting of three 2D arrays; each array represents the output of one particular scale.


The size of each array is  $n \times 85$ .

# YOLO Implementation

- `preds`, the output from YOLO model

Index ▲	Type	Size	Value
0	float32	(507, 85)	[[0.03275699 0.052097 0.3618731 ... 0. 0. 0. ...
1	float32	(2028, 85)	[[0.02476406 0.02638628 0.05257511 ... 0. 0. 0. ...
2	float32	(8112, 85)	[[0.00781204 0.00618281 0.01307961 ... 0. 0. 0. ...

Each row in the 2D array represents a box:



# YOLO Implementation

- Create 3 empty lists. One for storing the class of the object detected; one for the confidence level; last one for the position and the size of the box

```
> classId      = []  
> confidences  = []  
> boxes        = []
```

- Extract information from the output

```
> for scale in preds:           loop through each 2D array in the list  
    for pred in scale:         loop through each row in the 2D array  
        scores = pred[5:]      Extract all the class scores for this box  
        clss   = np.argmax(scores) Find out which class has the highest score  
        confidence = scores[clss] Extract the score of the class (which has the highest score)
```

```
> for scale in preds:
    for pred in scale:
        scores      = pred[5:]
        clss        = np.argmax(scores)
        confidence   = scores[clss]

        if confidence > 0.5:  Will consider the box only if the confidence is > 0.5
            xc        = int(pred[0]*imgWidth)      Get the actual box's center position, in x axis
            yc        = int(pred[1]*imgHeight)     Get the actual box's center position, in y axis
            w         = int(pred[2]*imgWidth)      Get the actual box's width
            h         = int(pred[3]*imgHeight)     Get the actual box's height
            x         = xc - w/2                  Get the actual box's top-left position, in x axis
            y         = yc - h/2                  Get the actual box's top-left position, in y axis

            classId.append(clss)
            confidences.append(float(confidence))   Must convert 'confidence' into float, else error in
            boxes.append([x, y, w, h])             later function (NMSBoxes)
```

Add the item and its corresponding parameters to the lists




# YOLO Implementation

- Perform non-maximal suppression to remove redundant overlapping bounding boxes and get a single box with the highest confidence score

```
> scoreThres      = 0.5
> nmsThres        = 0.4
> selected        = cv2.dnn.NMSBoxes(bboxes=boxes,
                                     scores=confidences,
                                     score_threshold=scoreThres,
                                     nms_threshold=nmsThres)
```

'selected' is m x 1 array,  
m is the number of  
objects selected



Source: <https://cloud.tencent.com/developer/article/1008757>

# YOLO Implementation

- Create colour set

```
> colorset = np.random.uniform(0,
                                255,
                                size=(len(classes),3))
```

- Extract colour, label for the detected class, and also retrieve the box position for plotting

```
> for j in selected[:,0]:  
    box      = boxes[j]  
    color    = colorset[classId[j]]  
    txtlbl   = str(classes[classId[j]])  
    x        = int(box[0])  
    y        = int(box[1])  
    w        = int(box[2])  
    h        = int(box[3])
```

because 'selected' is a m x 1 array, for convenience, we only take the first column

Get the colour for the detected class

Get the text label for the detected class

Convert values into int type to prevent potential TypeError when we draw rectangle or put text on image

# YOLO Implementation

- Continue previous for loop

```
> for j in selected:
```

```
    ...
```

```
    w      = round(box[2])
```

```
    h      = round(box[3])
```

```
    cv2.rectangle(img,  
                  (x,y),  
                  (x+w,y+h),  
                  color,  
                  2)
```

Draw the  
bounding box  
using  
rectangle

```
    cv2.putText(img,  
                txtlbl,  
                (x,y-5),  
                cv2.FONT_HERSHEY_SIMPLEX,  
                0.5,  
                color,  
                1,  
                cv2.LINE_AA)
```

Put up the txt  
label that  
specify the  
class



# YOLO implementation

- Note: the colour of the boxes will not be the same each time we run the code

